

Moving to RISC-V Vector

- A Practical Journey of
AI Operator Optimization

Guodong Xu,
RISCstar Solutions

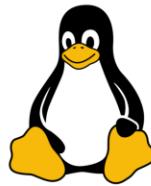


Who Am I ?

GUODONG XU

Chief Engineer,
Director of China Operations

RISCstar Solutions



Linux™



[*Op*]
[*en*]
[*BL*]
[*AS*]

Intel® ISA-L
Library

BLIS



open source project



arm



Agenda & Goals

Let's explore **RISC-V Vector (RVV)**'s practical application and performance by analyzing its implementation in two key open-source libraries: **OpenBLAS** (linear) and **SLEEF** (nonlinear).

- **Part 1: Analyze** the design logic for RVV in contrast to its Original Vector Code counterpart.
- **Part 2: Compare** RVV implementations side-by-side with code snippets.
- **Part 3: Demonstrate** vectorization speedup (Vector vs. Scalar) on RVV.
- **Part 4: Share** practical tooling insights using the *perf* utility.

Performance Measurement Environment & Toolchain

Software Platform



Fedora 42 Remix

Released by **Fedora-V Force team**

[Download Link: images.fedoravforce.org]

Hardware Platform



K1 SoC (RISC-V)

Manufactured by **SpacemIT**

RISCstar Toolchain



Built by experts with cross-architecture fluency, our toolchain is finely tuned for maximum performance on RISC-V Vector hardware.

Explore Now:

<https://riscstar.com/toolchain/>

Special thanks to **Mr. Haibin Liu** from the Fedora-V Force team, who provided invaluable support in conducting the performance measurement in OpenBLAS and using Perf.

PART I

RISC-V Vector: Considerations when Porting



Key Architectural Design Contrasts: RISC-V Vector vs. Source Platform

Vector Length Control

Source Platform (Reference)

Fixed at runtime (Hardware-Defined). Software discovers and adapts but cannot dynamically adjust.

RISC-V Vector (Advantage)

Dynamically Configurable: Software controls VL via the `vsetvl` instruction, allowing an optimal match between algorithm and hardware.

Register Grouping

Source Platform (Reference)

Fixed Register View: Lacks flexibility for logical grouping of vector registers.

Visual: Register Element Layout (e32m2 - 2 registers grouped)

VLEN=128: VLMAX=8, vl=8

Register 0 (4 element capacity):

0	1	2	3
---	---	---	---

Register 1 (4 element capacity):

4	5	6	7
---	---	---	---

Total: 8 elements (256 bits: 2 registers \times 128 bits)

VLEN=256 (K1): VLMAX=16, vl=8 (constrained by AVL)

Register 0 (8 element capacity):

0	1	2	3	4	5	6	7
---	---	---	---	---	---	---	---

Register 1 (8 element capacity):

.
---	---	---	---	---	---	---	---

3 elements processed (vl) | 8 available but not used

Total: 512 bits (2 registers \times 256 bits)

Understanding Application Vector Length (AVL)

RISC-V vectors are **VLEN-agnostic** through the `vsetvl` instruction, which takes an **Application Vector Length (AVL)** — the number of elements the software requests to process. The hardware returns the actual vector length (`vl`) based on this request and its capabilities.

How `vsetvl` Works:

Software Request
 $AVL = 8$
elements needed

`vsetvl` Calculation
 $VLMAX = (LMUL \times VLEN) / SEW$
 $vl = \min(AVL, VLMAX)$

Hardware Returns
 vl
actual length

- ✓ Vector operations process exactly `vl` elements
- ✓ Same code adapts to different VLEN implementations

Allows developers complete control over the execution length, which is essential for pursuing maximum manual performance optimization.

512 bits of data per operation (2x wider)

`LMUL=4 (4 registers grouped): 4x data per instruction`

<code>v0 (256 bits)</code>
<code>v1 (256 bits)</code>
<code>v2 (256 bits)</code>
<code>v3 (256 bits)</code>

1024 bits of data per operation (4x wider)

Code Comparison: VL-Agnostic Philosophy

Step 1: Loop Control

- **Loop control:** `svcntw()` vs returned `vl`.
- **Guarding work:**
predicate operand bundles (SVE) vs
masked vector results using explicit `vl` (RVV
masks only when needed).
- **Code impact:**
 - SVE keeps vector math “clean”, but demands predicated variants (`svmla_f32_m`);
 - RVV keeps operations non-predicated but requires explicit length plumbing (`__riscv_vfmacc_vv_f32m1`).

Source Platform / Original Vector

Code (predicate-driven)

- VL discovery is implicit. Loop advances using `svcntw()`.
- Per-iteration predicate `pg = svwhilelt_b32(i, n)` drives every instruction.
- Early exits (denormals, special cases)
use `svptest_any + svcompact` to segregate active lanes.

```
for (int i = 0; i < n; i += svcntw()) {  
    svbool_t pg = svwpg_b32(i, n);  
    svfloat32_t x = svld1(pg, &input[i]);  
  
    // polynomial under pg svmla_f32_m, etc.  
    svst1(pg, &output[i], acc);  
}
```

RISC-V RVV (length-driven)

- Loop front-loads `vl = __riscv_vsetvl_e32m1(n - i)`.
- `vl` controls both loop increment and each RVV intrinsic.
- Tail shrinks naturally via smaller `vl`; no predicates, but special handling uses `vmslt_vx_ + __riscv_vcompress`.

```
for (size_t i = 0; i < n; i += vl) {  
    size_t vl = __riscv_vsetvl_e32m1(n - i);  
    vfloat32m1_t x = __riscv_vle32_v_f32m1(&input[i], vl);  
  
    // polynomial using __riscv_vfmacc_vv_f32m1, etc.  
    __riscv_vse32_v_f32m1(&output[i], acc, vl);  
}
```

The core shift is from a **predicate-driven** model to an **explicitly length-driven** model.

Practical Guide: Key Porting Steps & Common Pitfalls

Step 2: Key Intrinsic & Concept Mapping

, where the list can be very long !

Recommended Resource for Deep Dive

For a comprehensive guide on intrinsic mapping and porting methodology, consider the upcoming **NEW COURSE: Porting Software to RISC-V (LFD 114)**. This training is developed to provide the cross-architecture fluency needed for real-world application porting.



Operation	Original Vector Intrinsics	RISC-V RVV Equivalent
Vector Load	<code>svld1(pg, &ptr[i])</code>	<code>__riscv_vle32_v_f32m1(&ptr[i], vl)</code>
FMA	<code>svmla_f32_m(pg, acc, x, h)</code>	<code>__riscv_vfmacc_vv_f32m1(acc, x, h, vl)</code>
Horizontal Sum	<code>float result = svaddv_f32(svtrue_b32(), acc_vec);</code>	<code>v_res = __riscv_vfredusum_vs_f32m8_f32m1(acc, v_res, vl);</code> <code>// Extract scalar result</code> <code>float result = __riscv_vfmv_f_s_f32m1_f32(v_res);</code>
Segmented Load / Structured Load	<code>svld3_u8(pg, base)</code> (up to 4 fields)	<code>vlseg3e8.v v8, (rs1), vm</code> (up to 8 fields)

PART II

AI Inference Operators

AI Inference Operators

Our analysis focuses on the core bottleneck of **LLM Inference**: the **Transformer Block**. Its workload can be broken down into two main operator categories:



Linear Operators (~90% of Compute)

Performance Attribution: Primarily dependent on high-performance linear algebra libraries like [OpenBLAS](#) / [BLIS](#).

Prefill Phase (Compute-Bound)

Dominated by **GEMM**: [seqlen, hidden_dim]

Decode Phase (Memory-Bound)

Dominated by **GEMV**: [1, hidden_dim]



Non-Linear Operators

Performance Attribution: Vectorization optimization primarily depends on high-performance math libraries like [SLEEF](#) / [LibM](#).

Key Operators

[GELU](#) [LayerNorm](#) [Softmax](#)

Mathematical Essence

[exp](#) [tanh](#) [sqrt](#) [Reduction/Sum](#)

PART III

Linear Operator Deep Dive (OpenBLAS)

$[Op]$ $[en]$
 $[BL]$ $[AS]$

The OpenBLAS RVV Strategy: "Fixed Small VL" for Cache Efficiency

The OpenBLAS implementation for RVV reveals a critical engineering trade-off that prioritizes cache efficiency for compute-bound tasks.

- **BLAS1/2 (e.g., GEMV, axpy):**
 - **Strategy:** Dynamic VL.
 - **Implementation:** Uses vsetvl to set VL to the number of remaining elements, making full use of the hardware's VLEN.
 - **Use Case:** Ideal for **Memory-Bound** operations to maximize memory bandwidth.
- **BLAS3 (MatMul Micro-kernels):**
 - **Strategy:** Intentionally uses a **Fixed Small VL** (e.g., VL=4 or 8).
 - **Implementation:** **DYNAMIC_ARCH** is used to compile specialized kernels for targets like ZVL128B and ZVL256B. Even on hardware with a wider VLEN, the kernel intentionally constrains VL.
 - **Engineering Goal:** To ensure the micro-kernel's tile size perfectly matches the L1 Cache, sacrificing VLEN utilization for maximum data locality. This is the key to performance in **Compute-Bound** tasks.

Code Comparison: SGEMM Core Computation (K-Loop Inner Body)

Source Platform Original Vector Code (predicate-driven)		RISC-V RVV (length-driven)
kernel/arm64/sgemm_kernel_sve_v1x8.S (874 lines)		kernel/riscv64/sgemm_kernel_8x8_zvl128b.c (791 lines)
Vector Load	<pre>; Software-pipelined, 2x unrolled K-loop .macro KERNELV1X8_M1 ld1w z1.s, p1/z, [pA] ; Load A[k+1] for next iter add pA, pA, lanes, lsl #2 ; Advance A pointer ; Compute with A[k] (z0), load B[k+1] for next iter fmla z16.s, p1/m, z0.s, z8.s ; acc0 += A[k] * B[k,0] ld1rw z8.s, p0/z, [pB] ; Broadcast B[k+1,0] fmla z17.s, p1/m, z0.s, z9.s ; acc1 += A[k] * B[k,1] ld1rw z9.s, p0/z, [pB, 4] ; Broadcast B[k+1,1] fmla z18.s, p1/m, z0.s, z10.s; acc2 += A[k] * B[k,2] ld1rw z10.s, p0/z, [pB, 8] ; Broadcast B[k+1,2] fmla z19.s, p1/m, z0.s, z11.s; acc3 += A[k] * B[k,3] ld1rw z11.s, p0/z, [pB, 12] ; Broadcast B[k+1,3] fmla z20.s, p1/m, z0.s, z12.s; acc4 += A[k] * B[k,4] prfm PLDL1KEEP, [pA, #A_PRE_SIZE] ; Prefetch A ld1rw z12.s, p0/z, [pB, 16] ; Broadcast B[k+1,4] fmla z21.s, p1/m, z0.s, z13.s; acc5 += A[k] * B[k,5] ld1rw z13.s, p0/z, [pB, 20] ; Broadcast B[k+1,5] fmla z22.s, p1/m, z0.s, z14.s; acc6 += A[k] * B[k,6] ld1rw z14.s, p0/z, [pB, 24] ; Broadcast B[k+1,6] fmla z23.s, p1/m, z0.s, z15.s; acc7 += A[k] * B[k,7] ld1rw z15.s, p0/z, [pB, 28] ; Broadcast B[k+1,7] add pB, pB, 32 ; Advance B pointer .endm</pre>	<pre>// Simple K-loop, compiler handles scheduling for (BLASLONG k = 1; k < K; k++) { // Load 8 scalars from B float B0 = B[bi + 0]; float B1 = B[bi + 1]; float B2 = B[bi + 2]; float B3 = B[bi + 3]; float B4 = B[bi + 4]; float B5 = B[bi + 5]; float B6 = B[bi + 6]; float B7 = B[bi + 7]; bi += 8; // Load vector from A (8 elements @ zvl128b) vfloat32m2_t A0 = __riscv_vle32_v_f32m2(&A[ai], gvl); ai += 8; // 8x vector-scalar FMA: acc += A * B (scalar broadcast implicit) result0 = __riscv_vfmacc_vf_f32m2(result0, B0, A0, gvl); result1 = __riscv_vfmacc_vf_f32m2(result1, B1, A0, gvl); result2 = __riscv_vfmacc_vf_f32m2(result2, B2, A0, gvl); result3 = __riscv_vfmacc_vf_f32m2(result3, B3, A0, gvl); result4 = __riscv_vfmacc_vf_f32m2(result4, B4, A0, gvl); result5 = __riscv_vfmacc_vf_f32m2(result5, B5, A0, gvl); result6 = __riscv_vfmacc_vf_f32m2(result6, B6, A0, gvl); result7 = __riscv_vfmacc_vf_f32m2(result7, B7, A0, gvl); }</pre>
Scalar Broadcast		
Vector-Scalar FMA		
Predication		

GEMM Performance Comparison & Attribution

For RISC-V, key observation points are to see what's the impact of the two:

1. VLEN (Vector Length)

Wider is Better: Yes or No?

2. LMUL (Vector Register Grouping, aka. Length Multiplier).

How much gains can we get when increasing LMUL from 1 to 2?

Note:

GENERIC means RISC-V Scalar

ZVL128B means RISC-V Vector 128bit VLEN

ZVL256B means RISC-V Vector 256bit VLEN

Peak Performance: Theoretical vs. Actual Speedup

Best performing configurations at matrix size 4096x4096

Gray = Theoretical Max | Color = Actual Performance

DGEMM (Double Precision)

Elements per vector: 64-bit → 2 (ZVL128B), 4 (ZVL256B)

GENERIC (Baseline)	1.0x
8176 MFLOPS	
ZVL128B (LMUL=4)	1.41x / 2.0x max
11552 MFLOPS	
ZVL256B (LMUL=1)	1.94x / 4.0x max
15857 MFLOPS	

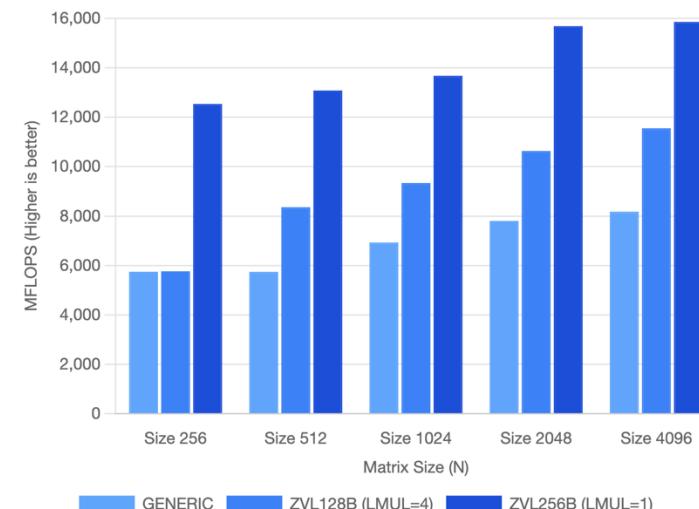
SGEMM (Single Precision)

Elements per vector: 32-bit → 4 (ZVL128B), 8 (ZVL256B)

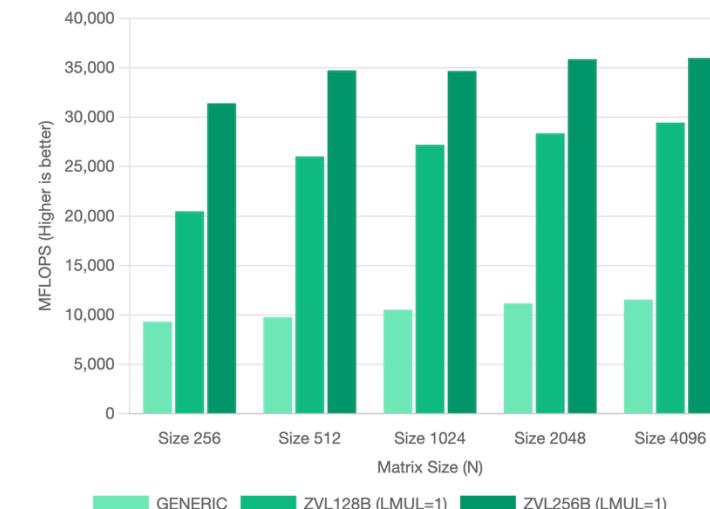
GENERIC (Baseline)	1.0x
11547 MFLOPS	
ZVL128B (LMUL=1)	2.55x / 4.0x max
29452 MFLOPS	
ZVL256B (LMUL=1)	3.12x / 8.0x max
35989 MFLOPS	

Raw Performance (MFLOPS vs. Matrix Si

DGEMM (Double Precision) - MFLOPS



SGEMM (Single Precision) - MFLOPS



For Reference, the Benchmark Testing Framework for OpenBLAS: <https://github.com/docularxu/OpenBLAS/tree/working.riscv64.optimizations>

Raw data and infographic published here: https://docularxu.github.io/OpenBLAS/webpages/rvv_benchmark_openblas_16Oct2025.html

GEMM Performance Comparison & Attribution

For RISC-V, key observation points are to see what's the impact of the two:

1. VLEN (Vector Length)

Wider is Better: Yes or No?

2. LMUL (Vector Register Grouping, aka. Length Multiplier).

How much gains can we get when increasing LMUL from 1 to 2?

Note:

GENERIC means RISC-V Scalar

ZVL128B means RISC-V Vector 128bit VLEN

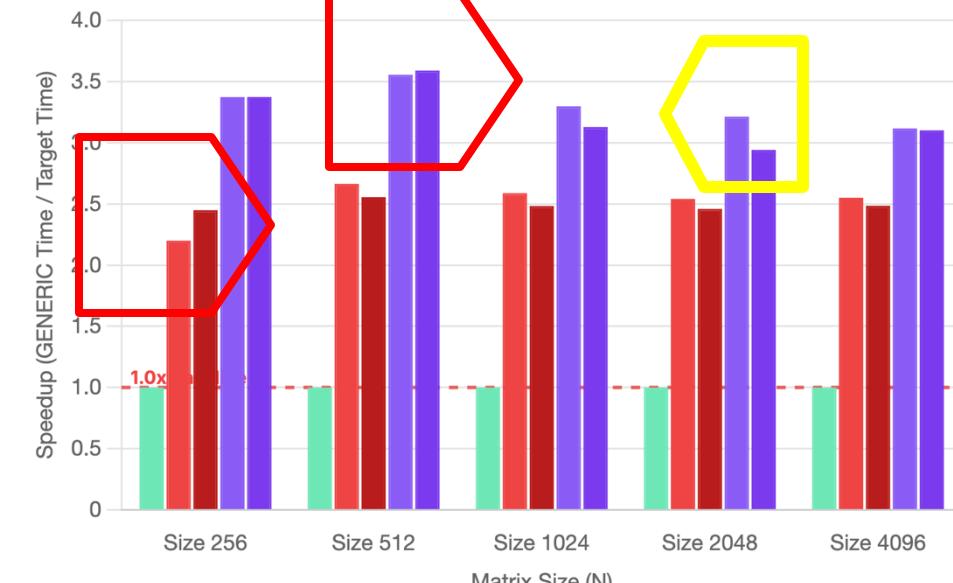
ZVL256B means RISC-V Vector 256bit VLEN

Speedup Analysis (vs. GENERIC Time)

DGEMM (Double Precision) - Speedup (GENERIC = 1.0x)



SGEMM (Single Precision) - Speedup (GENERIC = 1.0x)



GEMM Performance Comparison & Attribution

For RISC-V, key observation points are to see what's the impact of the two:

1. VLEN (Vector Length)

Wider is Better: Yes or No?

2. LMUL (Vector Register Grouping, aka. Length Multiplier).

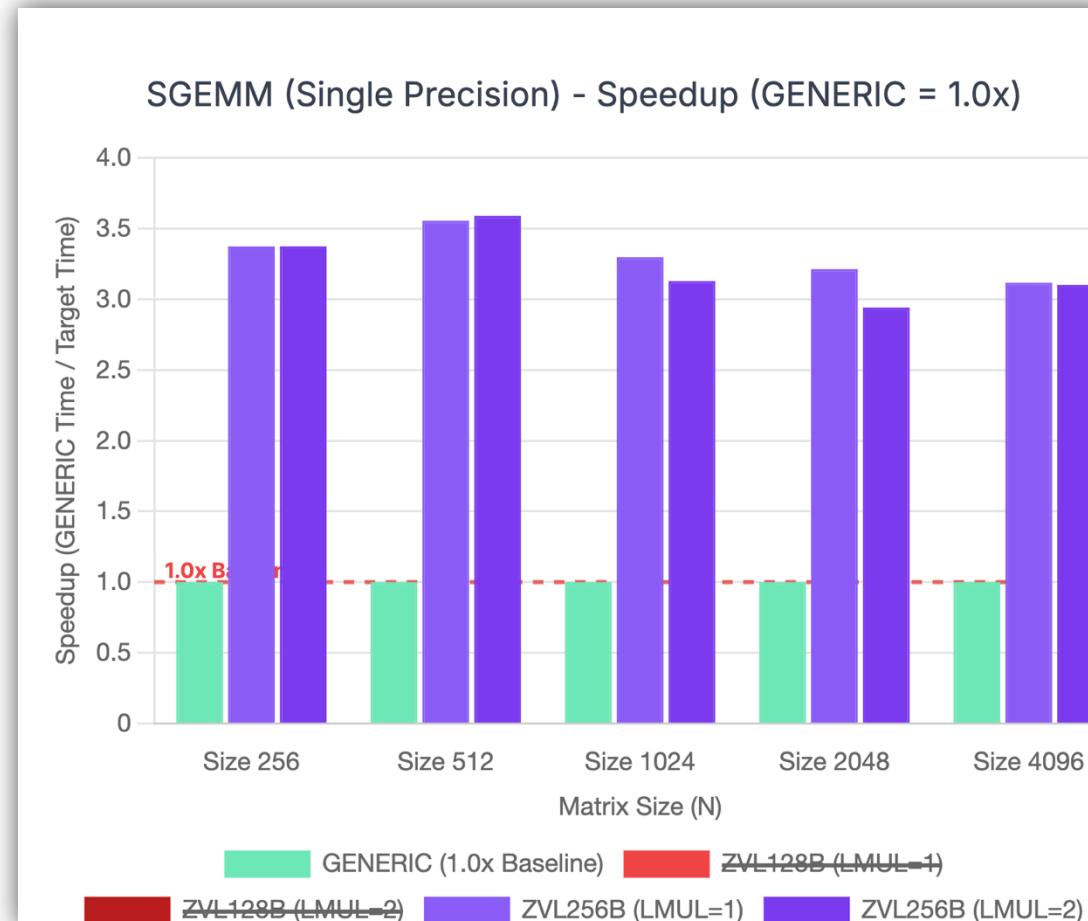
How much gains can we get when increasing LMUL from 1 to 2?

Note:

GENERIC means RISC-V Scalar

ZVL128B means RISC-V Vector 128bit VLEN

ZVL256B means RISC-V Vector 256bit VLEN



OBSERVATION:

For GEMM, LMUL provides minimal benefit and sometimes hurts performance

1. Only benefit: DGEMM on ZVL128B sees modest 8% improvement with LMUL=4, compensating for narrow 128-bit vectors
2. Matrix-size dependency: LMUL=2 shows 11% gain for SGEMM ZVL128B at 256×256 (22,805 vs 20,482 MFLOPS) but 3% degradation at 4096×4096 (28,696 vs 29,451 MFLOPS)
3. Recommendation: For production deployments on K1, use matrix-size adaptive LMUL tuning: LMUL=2 for small matrices (<1024), LMUL=1 for large matrices (≥ 1024), except DGEMM ZVL128B where LMUL=4 provides 8% gain

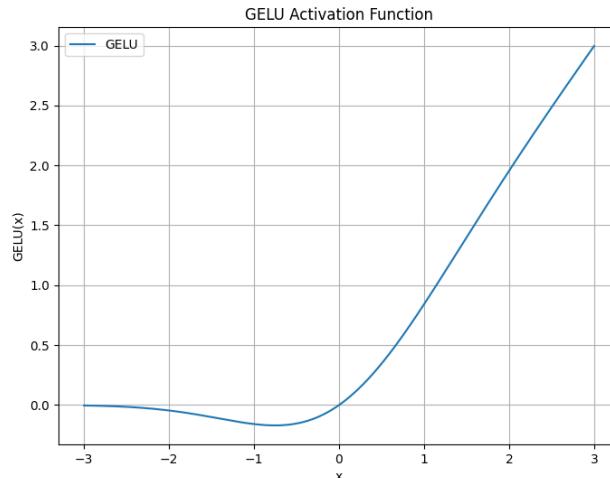
PART IV

Non-Linear Operator Deep Dive (SLEEF)



Introduction to Non-Linear Operators

Non-linear operators are essential for LLM performance, but they rely heavily on optimized standard mathematical functions ([libm](#)).



Softmax

Converts raw scores (logits) to probabilities.

Core Dependencies:

- `exp` (Exponential function)
- `reduction / sum`

GeLU (Gaussian Error Linear Unit)

A standard activation function in Transformer models.

Core Dependencies:

- `exp` (Often via error function approximation)
- `tanh` (Hyperbolic Tangent)

LayerNorm (Layer Normalization)

Stabilizes the learning process across feature dimensions.

Core Dependencies:

- `sqrt` (Square Root)
- `reduction / sum`

- `exp`: (Used in Softmax, GeLU)
- `tanh`: (Used in GeLU)
- `sqrt`: (Used in LayerNorm)
- `reduction / sum`: (Used in LayerNorm, Softmax)

$$e^x$$

Exponential (exp)

Widely used in **Softmax** and **GeLU** activation functions. Its vectorization is a key performance bottleneck for these critical AI operators.

$$\tanh$$

Hyperbolic Tangent (tanh)

Another core component of the **GeLU** activation function, crucial for the model's non-linear expression capabilities.

$$\sqrt{x}$$

Square Root (sqrt)

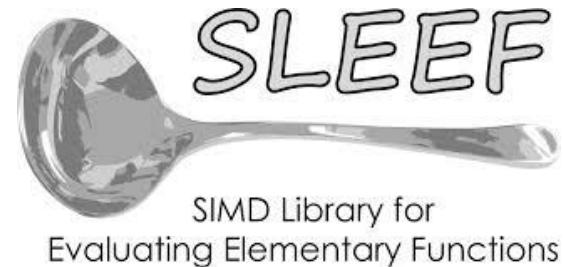
Primarily used for variance calculation in **LayerNorm**, a key step in maintaining training stability.

Vectorizing AI Math Operators

Why SLEEF?

The performance of non-linear operators like *GeLU*, *Softmax*, and *LayerNorm* depends on the efficient vectorization of fundamental math functions.

- **Core Challenge:** How to apply SIMD principles to complex functions like *exp*, *tanh*, and *sqrt*.
- **The Solution:** High-performance math libraries like *SLEEF* provide hand-tuned vector implementations of these functions for different architectures.
- **Our Focus:** Comparing how SLEEF leverages the unique features of RISC-V Vector and other VLA platforms to accelerate these critical AI operators.



SLEEF Performance Measured (Vector vs. Scalar) Speedup Comparison

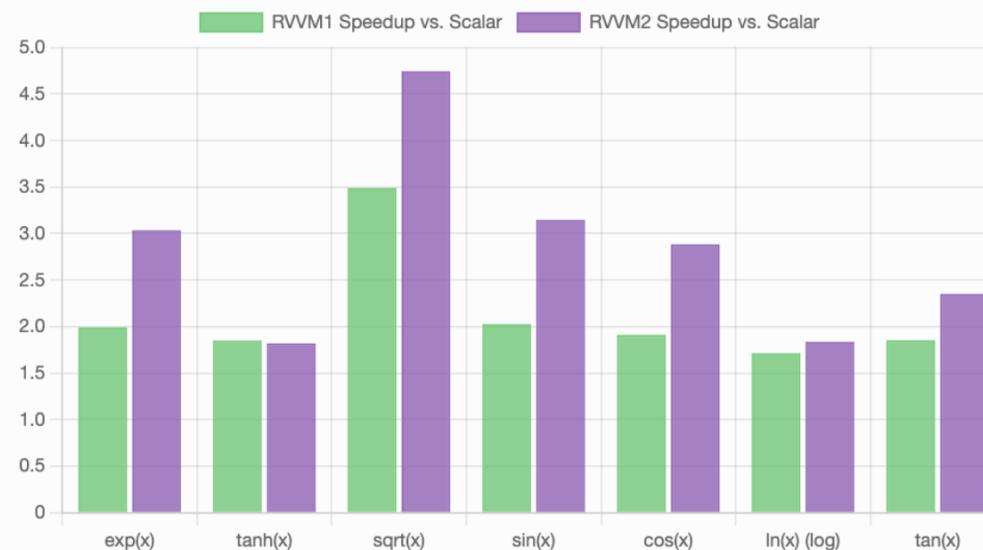
Three categories are measured:

- Scalar
- RVV with LMUL=1
- RVV with LMUL=2

LMUL = Length Multiplier

Double Precision Speedup (vs. Sleef Scalar)

Vector Speedup (Double Precision)



Speedup factors of RVVM1 and RVVM2 compared to the Sleef Scalar implementation for double precision functions. Higher bars indicate greater speedup.

Single Precision Speedup (vs. Sleef Scalar)

Vector Speedup (Single Precision)



Speedup factors of RVVM1 and RVVM2 compared to the Sleef Scalar implementation for single precision functions. Higher bars indicate greater speedup.

SLEEF Performance Measured (LMUL = 2 vs. LMUL = 1) RVVM2 Gain Over RVVM1

RISC-V provides a powerful feature not found in SVE: **Vector Register Grouping (LMUL, ie. Length Multiplier)**. This allows software to logically combine multiple registers to operate on wider data or reduce loop overhead.

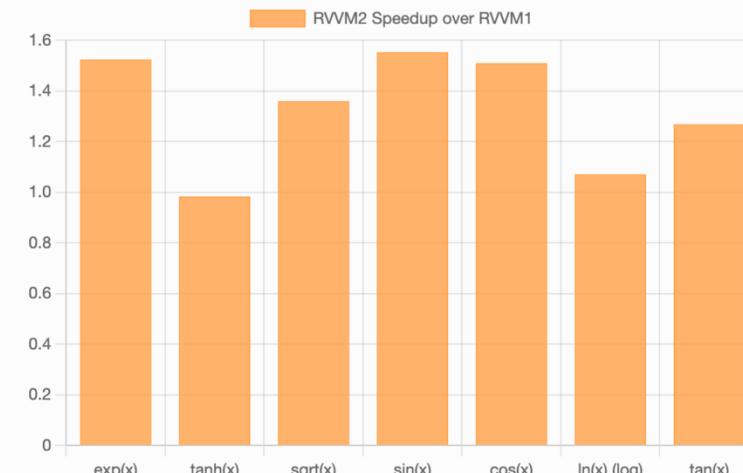
Three categories are measured:

- Scalar
- RVV with LMUL=1
- RVV with LMUL=2

SLEEF Config	Description	Relative Performance
RVVM1 (LMUL=1)	Standard, one register per operation.	1.0x (Baseline)
RVVM2 (LMUL=2)	Groups two registers	~1.4x <i>Caution: tanh(x) sees no gains when growing LMUL</i>

RVVM2 Speedup vs. RVVM1 (Double Precision)

RVVM2 Speedup vs. RVVM1 (Double Precision)



Speedup factor of RVVM2 using RVVM1 as the baseline. This highlights the performance benefit of increasing LMUL from 1 to 2.

RVVM2 Speedup vs. RVVM1 (Single Precision)

RVVM2 Speedup vs. RVVM1 (Single Precision)



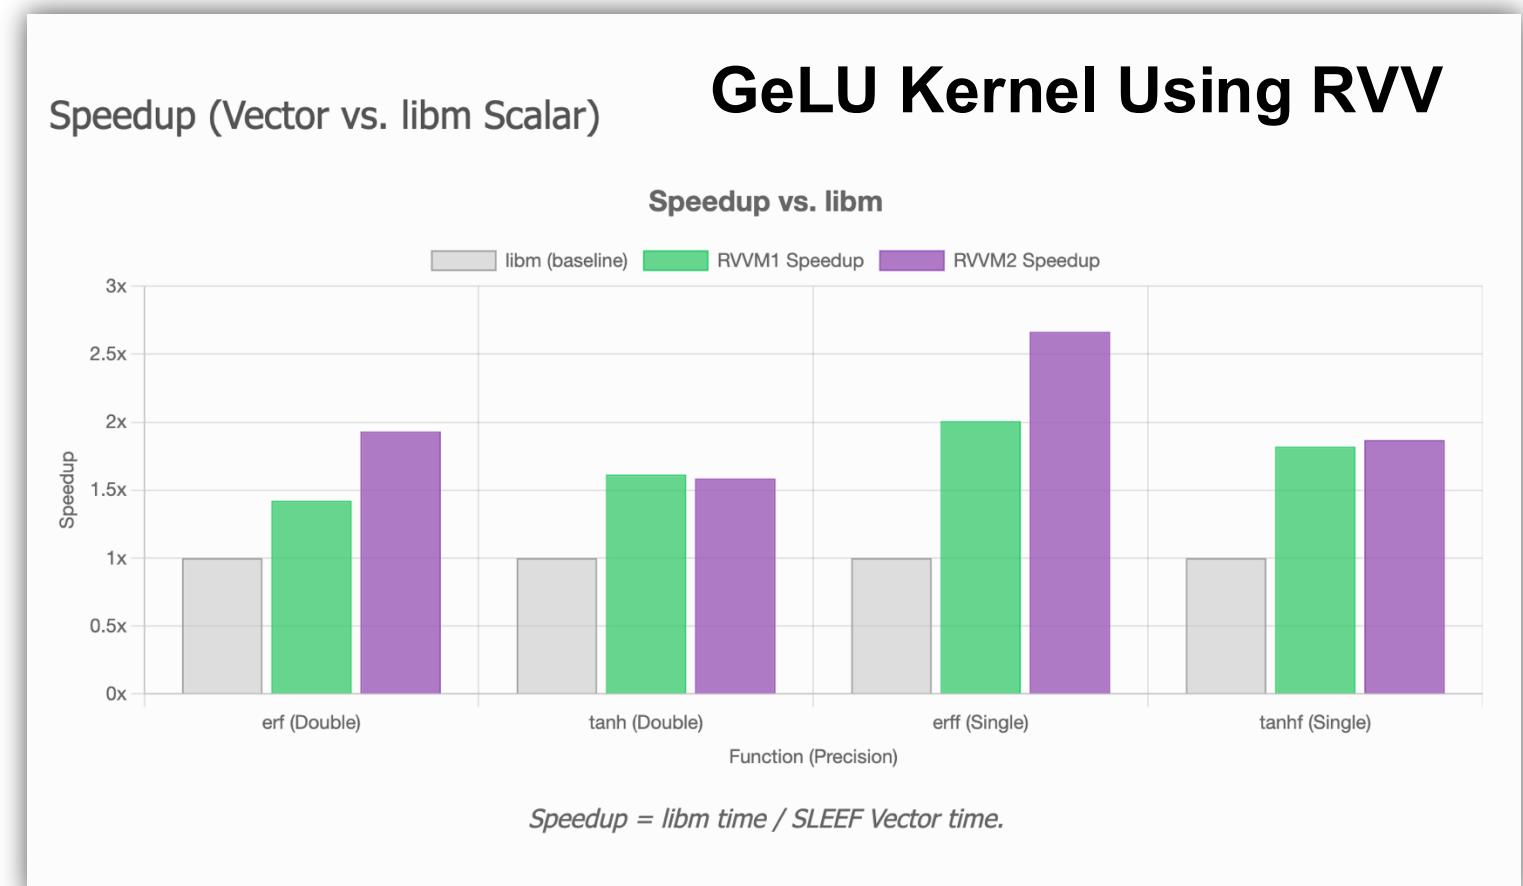
Speedup factor of RVVM2 using RVVM1 as the baseline. This highlights the performance benefit of increasing LMUL from 1 to 2.

SLEEF Performance Measured (LMUL = 2 vs. LMUL = 1) GeLU on RISC-V

Comparison of libm scalar vs. SLEEF vector implementations in GeLU.

- Scalar
- RVV with LMUL=1
- RVV with LMUL=2

SLEEF Configuration	Relative Performance
Libm (scalar)	1.0x (Baseline)
RVVM1 (LMUL=1)	1.5x speedup
RVVM2 (LMUL=2)	~2.0x speedup



PART V

Tools and Toolchains

Tooling & Observability

Performance Monitoring (perf):

- RISC-V Platform (K1): Support is available in their vendor tree. Refer to:
<https://bianbu.spacemit.com/en/development/perf/>
- The SpacemIT K1 provides some essential events, but the ecosystem is still evolving. Not upstreamed.
- Improving perf support is a key area for community contribution.

SpacemIT K1/M1 Supported Perf Events

Event	Status	Event	Status	Event	Status	Event	Status
cycles	✓	LLC-prefetches	✗	dtlb_load_miss	✓	ecall_inst	✓
branches	✓	LLC-prefetch-misses	✗	dtlb_store_miss	✓	failed_sc_inst	✓
branch-misses	✓	dTLB-loads	✗	itlb_load_miss	✓	fence_inst	✓
instructions	✓	dTLB-load-misses	✓	jtlb_miss	✓	fp_div_inst	✓
bus-cycles	✗	dTLB-stores	✗	i1d_access	✓	fp_inst	✓
ref-cycles	✗	dTLB-store-misses	✓	i1d_amr_active	✓	fp_load_inst	✓
page-faults	✓	dTLB-prefetches	✗	i1d_excl_evict	✓	fp_store_inst	✓
context-switches	✓	dTLB-prefetch-misses	✗	i1d_load_access	✓	load_inst	✓
cpu-migrations	✓	iTLB-loads	✗	i1d_load_miss	✓	lr_inst	✓
stalled-cycles-frontend	✓	iTLB-load-misses	✓	i1d_miss	✓	mult_inst	✓
stalled-cycles-backend	✓	branch-loads	✗	i1d_prefetch_hit	✓	sc_inst	✓
task-clock	✓	branch-load-misses	✗	i1d_prefetch_refill	✓	store_inst	✓
cache-misses	✗	node-loads	N/A	i1d_store_access	✓	unaligned_load_inst	✓
cache-references	✗	node-load-misses	N/A	i1d_store_miss	✓	unaligned_store_inst	✓

Examples:

(Right) K1 proprietary/vendor-specific events
vs.

(Left) the Linux common events

Summary & Call to Action

Key Takeaways from Our Porting Journey:

- 1. Master VL Control:** For compute-bound tasks, **data locality is king**. Don't chase maximum LMUL. Use RVV's vsetvl to manually craft small VL loops that fit the L1 cache, just as OpenBLAS does.
- 2. Leverage the Math Libraries:** Don't reinvent the wheel for non-linear functions. **Use and contribute to libraries like SLEEF**. Improving these libraries is the fastest path to unlocking RVV's performance for AI workloads.
- 3. Embrace Observability:** Trust, but verify. Use perf and analyze the generated assembly to ensure your VL settings, tail policies, and memory accesses are what you expect.
- 4. Engage with the Ecosystem:** RISC-V's greatest strength is its community. Submit patches and provide feedback to upstream toolchains like LLVM and GCC. Every improvement benefits everyone.



Optimize with the RISCstar Toolchain

Built by experts with cross-architecture fluency, our toolchain is finely tuned for maximum performance on RISC-V Vector hardware.

Explore Now:

<https://riscstar.com/toolchain/>



Connect with me:

GUODONG XU

EMAIL: guodong@riscstar.com

LINKEDIN: <https://www.linkedin.com/in/docularxu/>

GITHUB: <https://www.github.org/docularxu>

THANK YOU!



SLEEF Performance Measured (Vector vs. Libm reference) Speedup Comparison

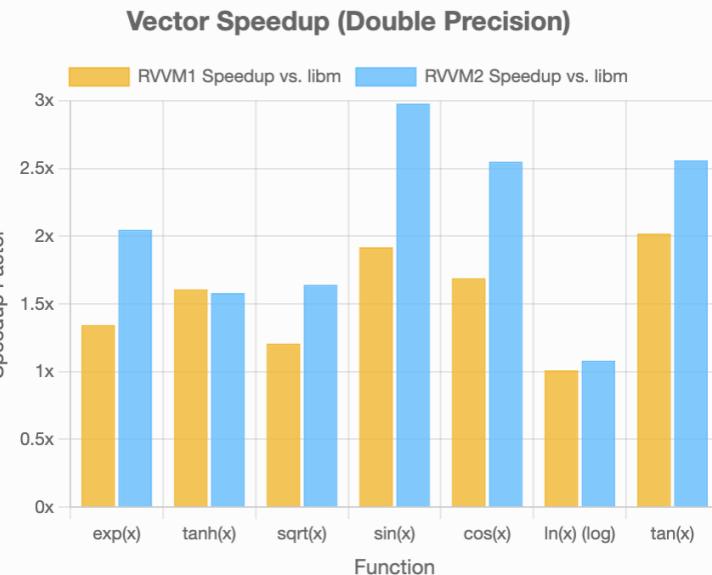
Speedup Comparison (Vector vs. libm Reference)

Three categories are measured:

- Libm Reference
- RVV with LMUL=1
- RVV with LMUL=2

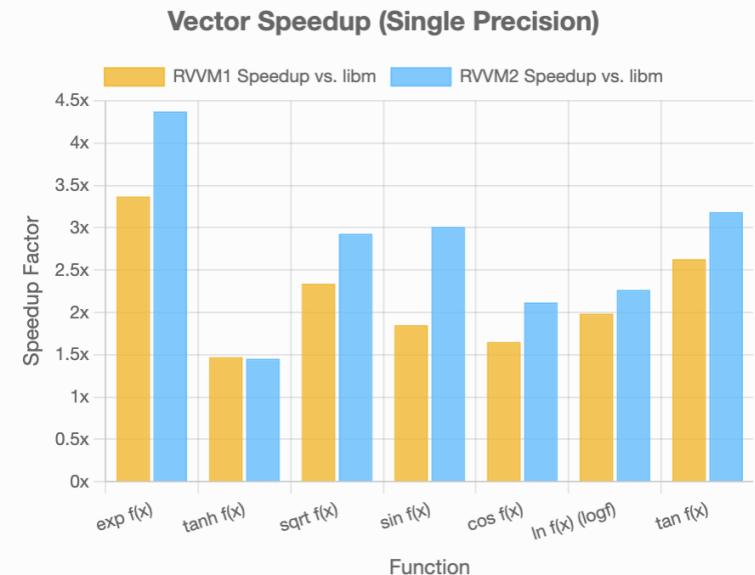
LMUL aka. Length Multiplier

Double Precision Speedup (vs. libm Reference)



Speedup factors of RVVM1 and RVVM2 compared to the libm reference implementation. This shows the vector gain over standard system libraries.

Single Precision Speedup (vs. libm Reference)



Speedup factors of RVVM1 and RVVM2 compared to the libm reference implementation. This shows the vector gain over standard system libraries.